

xkcd #1831
(slightly edited)

caches


CPEN212 2022 W2



a place of mind

THE UNIVERSITY OF BRITISH COLUMBIA

Q. what is a cache?



key idea: reuse

temporal reuse

- **observation:** some data often accessed **repeatedly**
 - e.g., loop counter + other variables accessed in loop
- **idea:** keep frequently accessed data nearby
- **cache:** small, fast storage near use location
 - e.g., on-chip memory near CPU = cache for off-chip DRAM
 - e.g., physical DRAM = cache for on-disk virtual memory
 - e.g., proxy web server / CDN = cache for remote webpages

spatial reuse

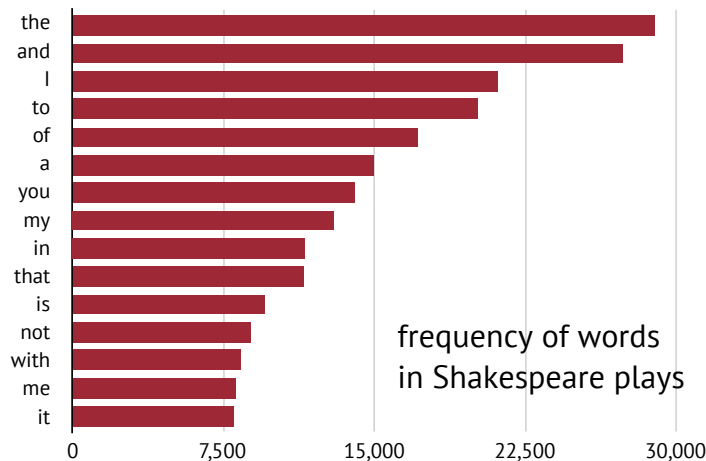
- **observation:** data often accessed **in local clusters**
 - e.g., iterating through an array
- **idea:** also cache data **spatially near** recent accesses
- **cache line:** a range of data brought in cache at once
 - e.g., several contiguous addresses (e.g., 64–128 bytes)
 - e.g., nearby pixels (in a GPU texture cache)

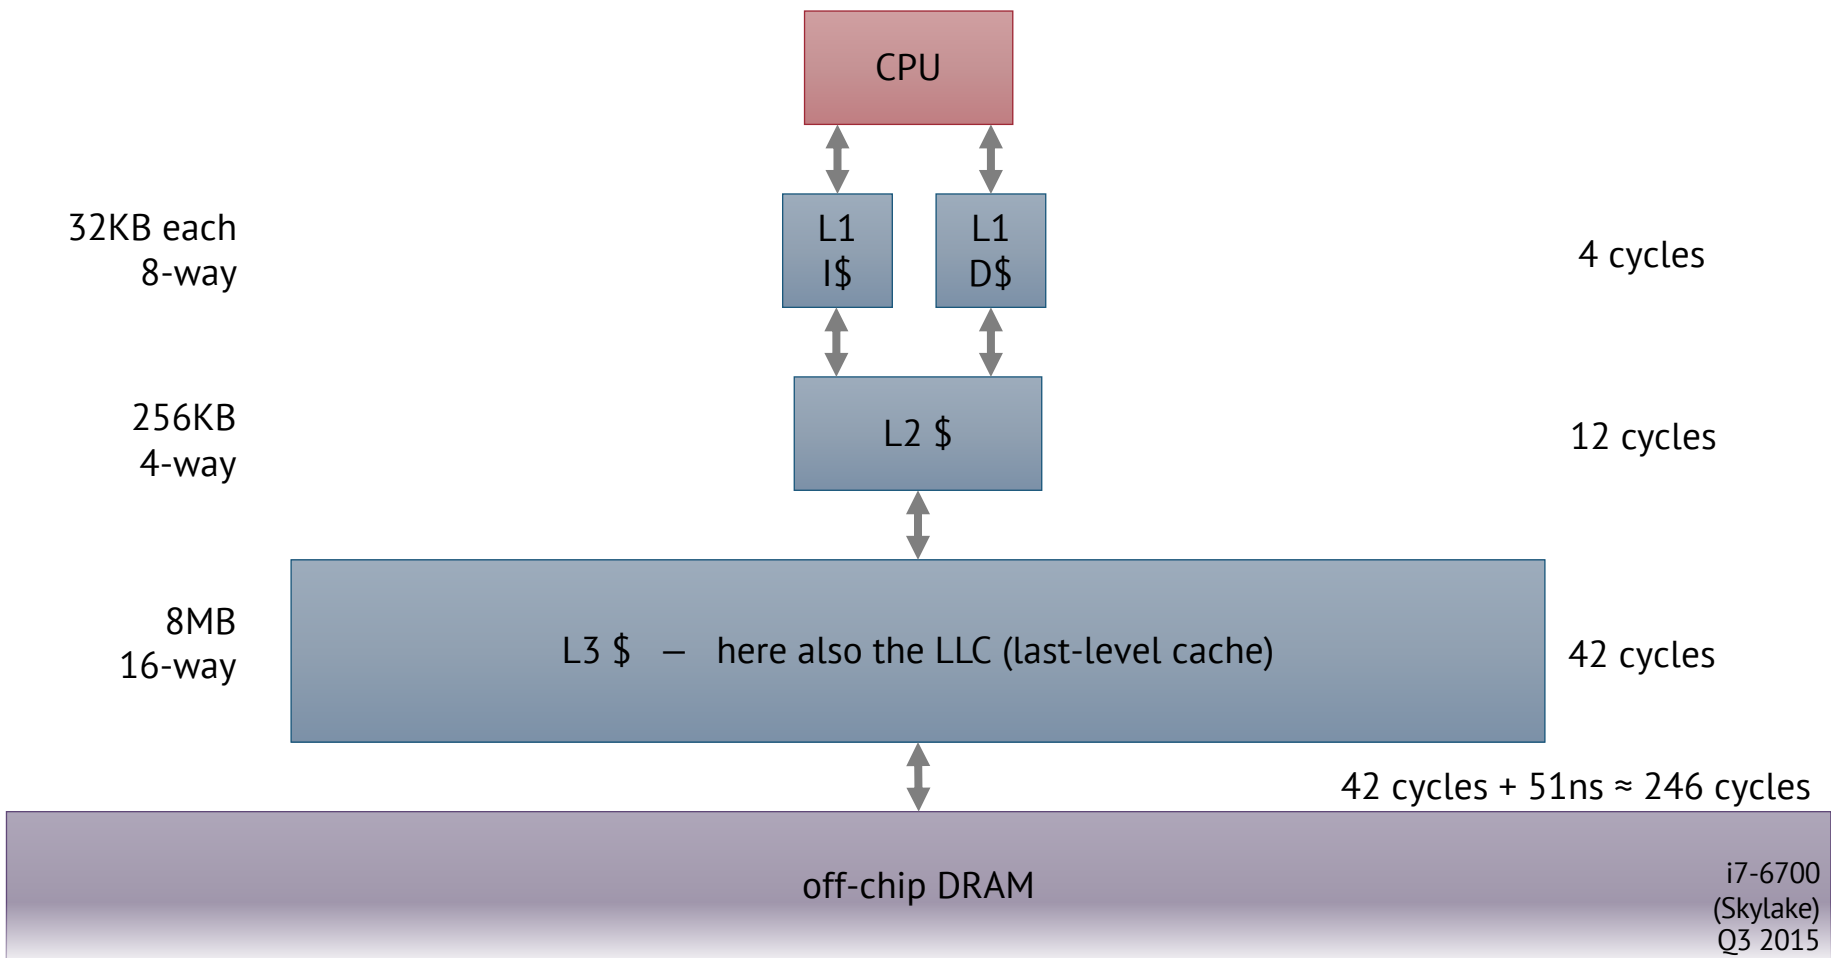
Q. what kind of reuse was the cat's wallet?

key idea: hierarchy

cache hierarchies

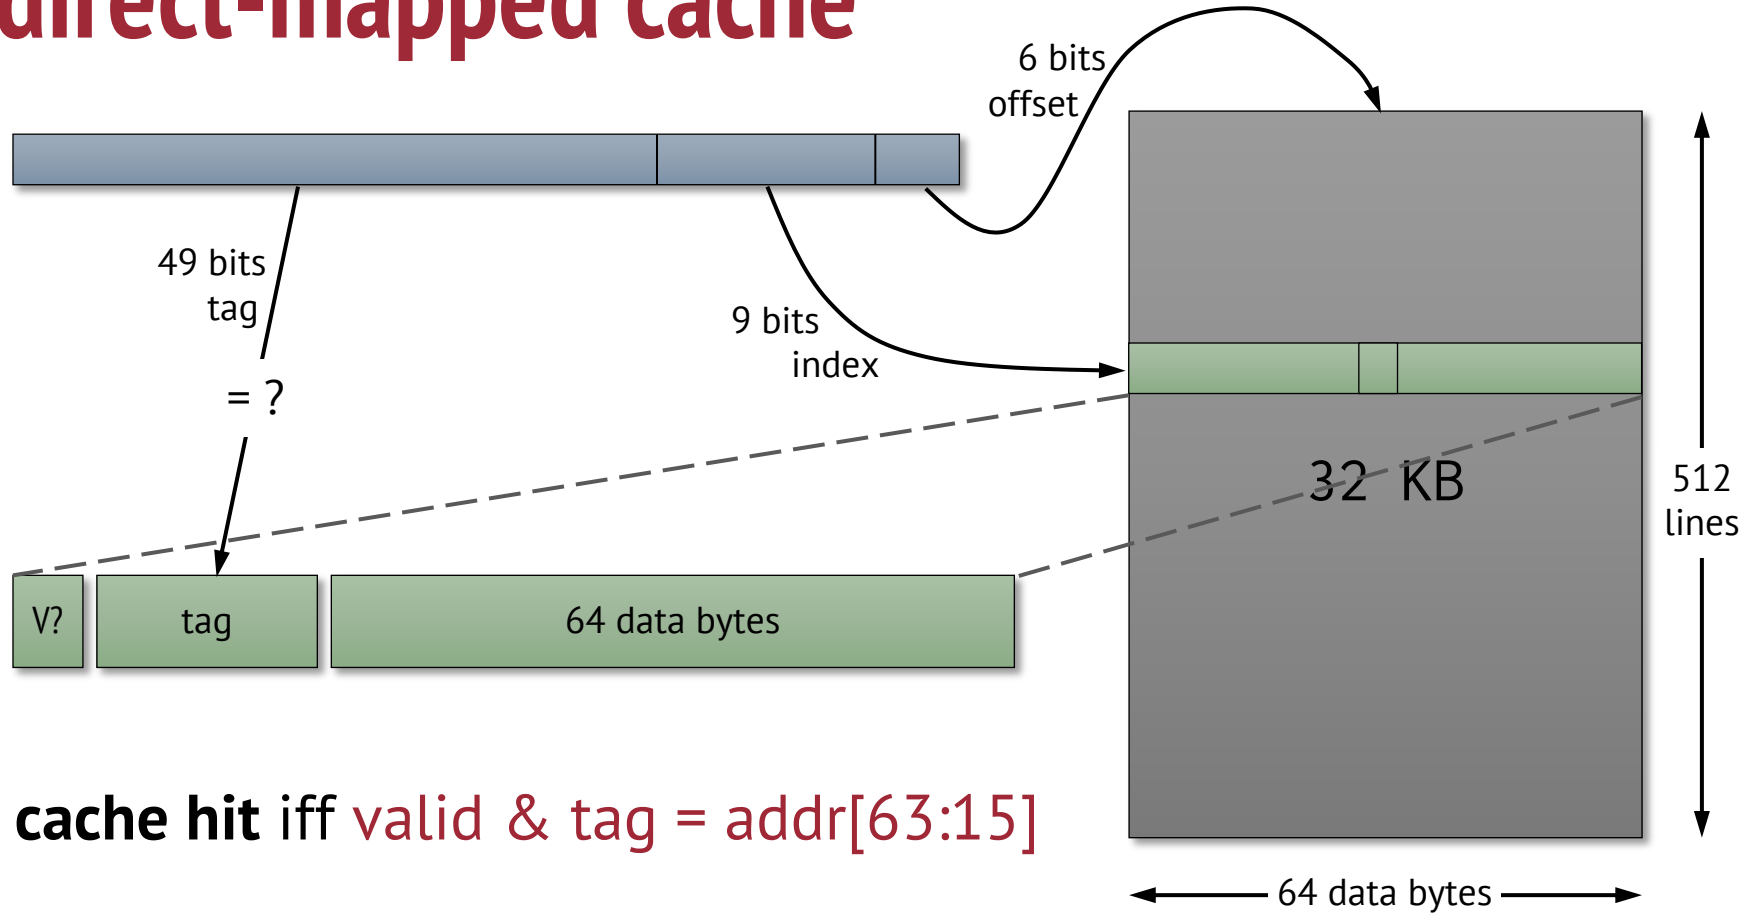
- memory latency $\sim \sqrt{\text{size}}$
(w/ same technology)
- few data accessed very often
 - e.g., words in language
- **idea:** a fast small \$ in front of a larger, slower \$
 - and maybe another, even larger, even slower \$...





cache organization

direct-mapped cache




direct-mapped cache problems



- same index bits → **conflict** → eviction
- if indices random, $\text{Pr}[\text{conflict}] > 50\%$ if only **27 lines**
- even worse: **pathological patterns**
 - e.g., if accessing many words 4096 bytes apart
 - **only 8** cache lines used even if the cache has 512!
 - all accesses past the 8th **evict** some prior cache line

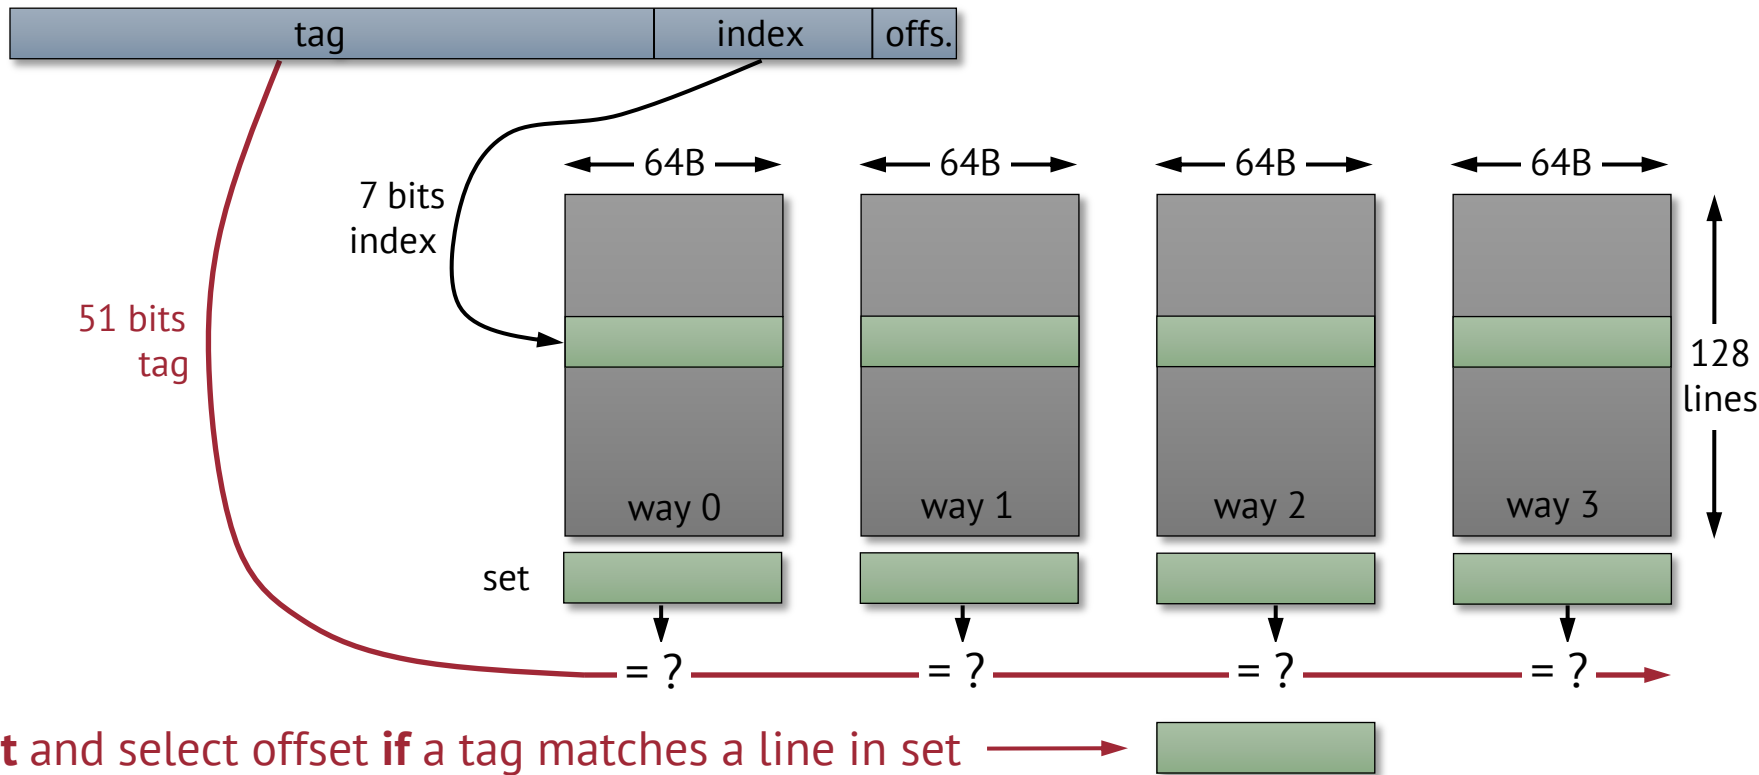
fully associative cache



compromise: **set-associative caches**

- **idea**: organize cache lines in **sets** (e.g., 4 lines / set)
- use index bits to **select sets**
- **search** for tag in all **ways** in the selected set
 - some extra hardware but $\# \text{ ways} \ll \# \text{ indices}$ so not bad
- now conflict harder: need **$|\text{set}|+1$** equal indices

set-associative cache



replacement policies

- if want to insert but set full, which line to evict?
- ideal replacement policy (Bélády MIN algorithm):
evict line re-referenced **furthest** in the future
- but, can't know the future T_T
- usually **least-recently used** (LRU) or **random**

writing cache-friendly code

using caches efficiently

- **lay out** data accessed together inside one cache line
 - after the first miss the rest of cache line will hit
 - Q: **spatial or temporal locality?**
- reuse data in chunks **that fit in the cache** (tiles)
 - data likely to still be in cache when reused
 - Q: **spatial or temporal locality?**
- focus on **innermost cache first (L1)**, then outer caches

finding reuse: **matrix multiplication**

key idea: find matrix indices **constant** across a loop

need to loop through r, c

need to loop through i


$$\forall r, c: z[r, c] \leftarrow \sum_i x[r, i] \cdot y[i, c]$$

z **not indexed** with i
→ z **reused** for every i
(for accumulation)

x **not indexed** with c
→ x **reused** for every c

y **not indexed** with r
→ y **reused** for every r

let's check: **is y reused for every r?**

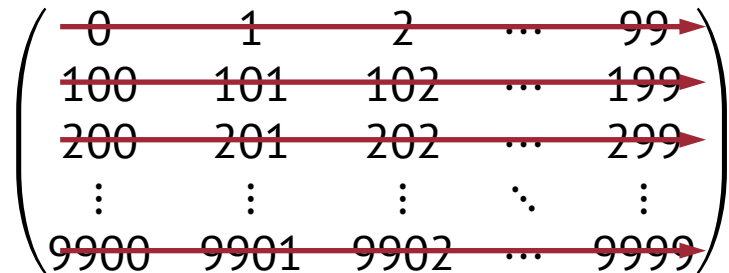


verify x and z reuse yourself (same logic)

Q. spatial or temporal reuse?

interlude: 2D matrix representations

row-major order

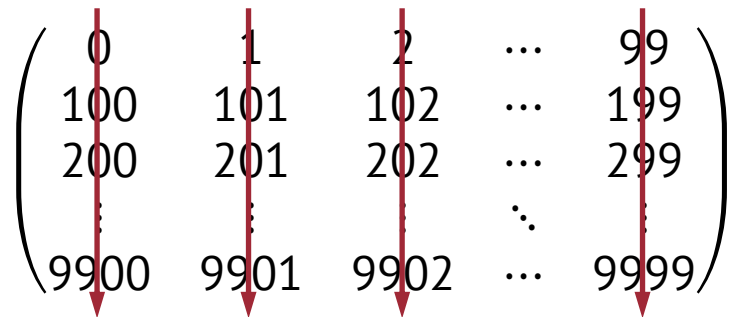
$$\begin{pmatrix} 0 & 1 & 2 & \dots & 99 \\ 100 & 101 & 102 & \dots & 199 \\ 200 & 201 & 202 & \dots & 299 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 9900 & 9901 & 9902 & \dots & 9999 \end{pmatrix}$$


```
double *matrix = { 0, 1, 2, ..., 99, 100, 101, 102, ..., 9999 };
```

```
double element = matrix[100 * row + col];
```

interlude: 2D matrix representations

column-major order

$$\begin{pmatrix} 0 & 1 & 2 & \dots & 99 \\ 100 & 101 & 102 & \dots & 199 \\ 200 & 201 & 202 & \dots & 299 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 9900 & 9901 & 9902 & \dots & 9999 \end{pmatrix}$$


```
double *matrix = { 0, 100, 200, ..., 9900, 1, 101, 201, ..., 9999 };
```

```
double element = matrix[row + 100 * col];
```

example: matrix multiplication

- problem setup: $z = x \cdot y$
 - x, y, z are 64×64 matrices of doubles (8B each)
 - 16KB cache, 4-way set-associative
- Q. does each matrix fit in the cache?
- Q. does one row or column fit in the cache?
- plan: maximize reuse at **row / column** level

```


N = 64; // 3 × 4096-elt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
  for (c = 0; c < N; ++c) // output col
    for (i = 0; i < N; ++i) // x col, y row
      z[N * r + c] += x[N * r + i] * y[N * i + c];

```

read *and* write

read

read



Q. expected reuse?

Q. cache miss rates?

```
$ valgrind --tool=cachegrind -D1=16384,4,64 ...
```

...

```
D   refs:          1,056,375 (792,821 rd + 263,554 wr)
```

```
D1  misses:         268,094 (267,960 rd +      134 wr)
```

```
D1  miss rate:       25.4% ( 33.8% + 0.1% )
```



~1 in 3 reads miss???

most writes hit


```

N = 64; // 3 × 4096-elt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
    for (c = 0; c < N; ++c) // output col
        for (i = 0; i < N; ++i) // x col, y row
            z[N * r + c] += x[N * r + i] * y[N * i + c];



```



```

N = 64; // 3 × 4096-elt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
    for (c = 0; c < N; ++c) // output col
        for (i = 0; i < N; ++i) // x col, y row
            z[N * r + c] += x[N * r + i] * y[N * i + c];

```




problem: all N^3 accesses to y miss in the cache!

```


N = 64; // 3 × 4096-elmt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
  for (c = 0; c < N; ++c) // output col
    for (i = 0; i < N; ++i) // x col, y row
      z[N * r + c] += x[N * r + i] * y[N * i + c];

```



observation: want i loop to be **outside** of c loop

idea:
innermost
loop swap





Q. does swapping maintain functional correctness?

```

N = 64; // 3 × 4096-elmt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
    for (i = 0; i < N; ++i) // x col, y row
        for (c = 0; c < N; ++c) // output col
            z[N * r + c] += x[N * r + i] * y[N * i + c];

```





Q. what happens when i increments?

```

N = 64; // 3 × 4096-elt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
    for (i = 0; i < N; ++i) // x col, y row
        for (c = 0; c < N; ++c) // output col
            z[N * r + c] += x[N * r + i] * y[N * i + c];

```





Q. estimated read / write miss rates for large N?

```

N = 64; // 3 × 4096-elt matrices (32KB each); 16KB cache
for (r = 0; r < N; ++r) // output row
    for (i = 0; i < N; ++i) // x col, y row
        for (c = 0; c < N; ++c) // output col
            z[N * r + c] += x[N * r + i] * y[N * i + c];

```




D1 miss rate: 3.2% (4.3% + 0.1%)

processing in cache-sized chunks

- loop reordering \rightarrow better reuse
- but total footprint still **too big for the cache**

- idea: **tiling**
 - split matrix into **submatrices**
 - as big as we can where all three still **fit in the cache**
 - finish tile before moving on



$N = 64$; $T = 16$;

```
for (r0 = 0; r0 < N; r0 += T) // output row tile
```

```
    for (c0 = 0; c0 < N; c0 += T) // output col tile
```

```
        for (i0 = 0; i0 < N; i0 += T) // x col, y row tile
```

```
            for (r = r0; (r < r0 + T) && (r < N); ++r) // output row
```


```
                for (i = i0; (i < i0 + T) && i < N; ++i) // x col, y row
```

```
                    for (c = c0; (c < c0 + T) && (c < N); ++c) // output col
```

```
                         $z[N * r + c] += x[N * r + i] * y[N * i + c];$ 
```

} loop through tiles


} loop through
elements
inside tile



L1 + L2 cache?

- **two** levels of cache
→ **two** tiling levels
- reuse T1 in L1\$
(inner tile)
- reuse T2 in L2\$
(outer tile)


Q. what about L3\$?

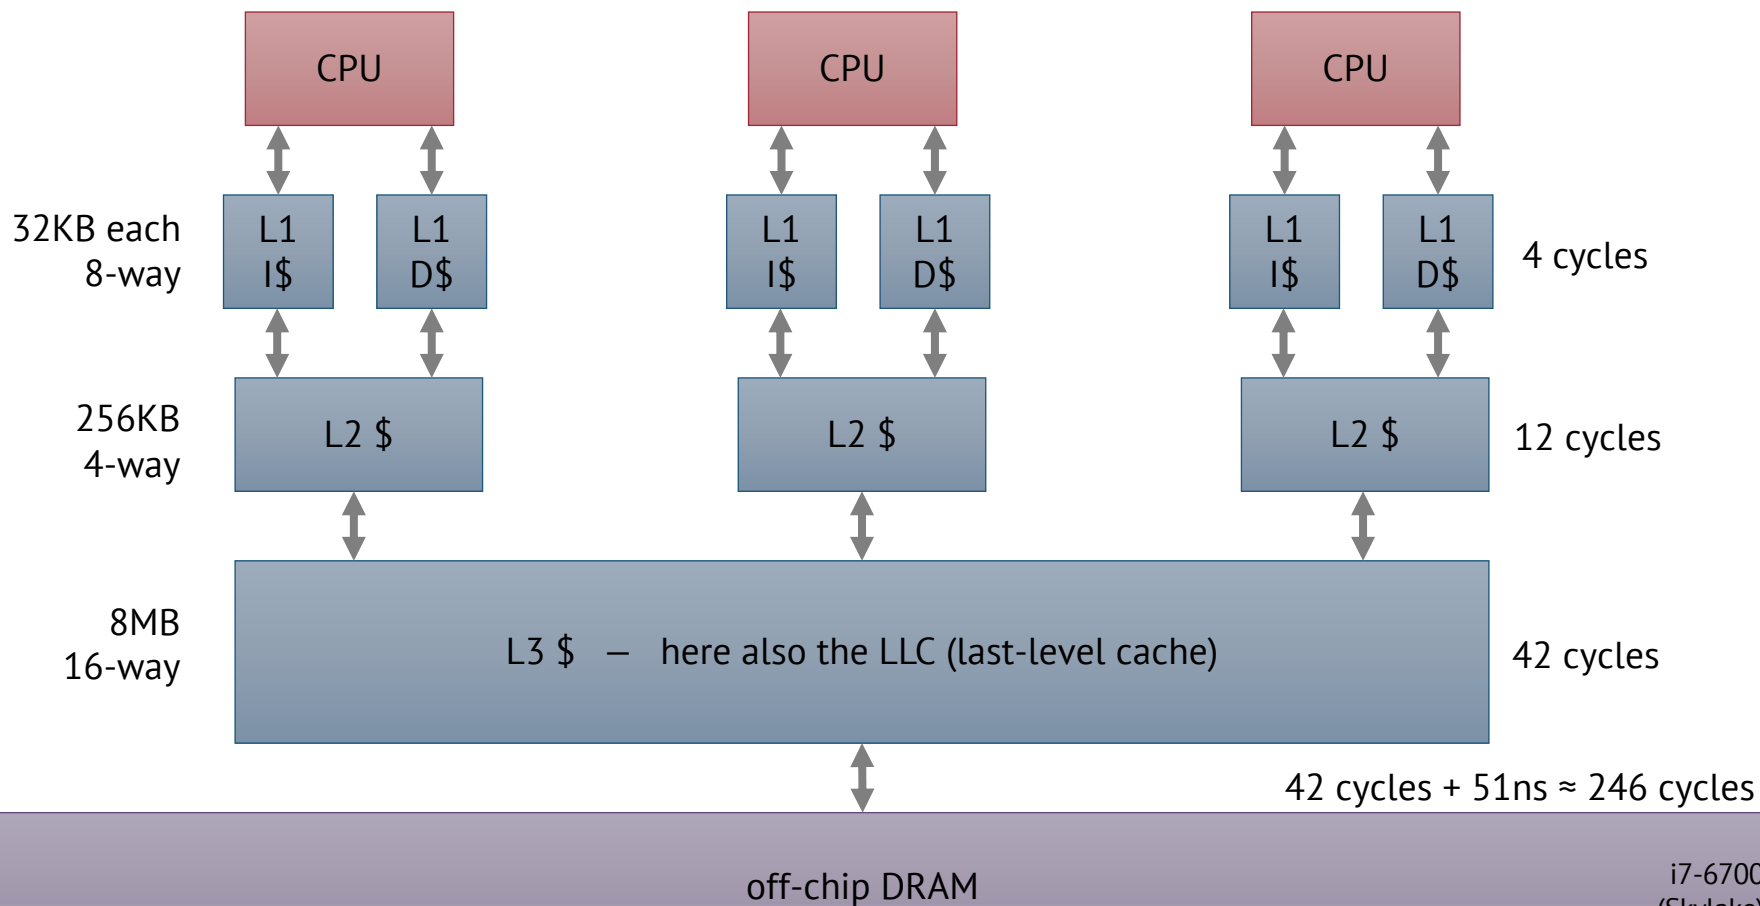


L1 + L2 cache?

- **two** levels of cache
→ **two** tiling levels
- reuse T1 in L1\$
(inner tile)
- reuse T2 in L2\$
(outer tile)

Q. what about L3\$?



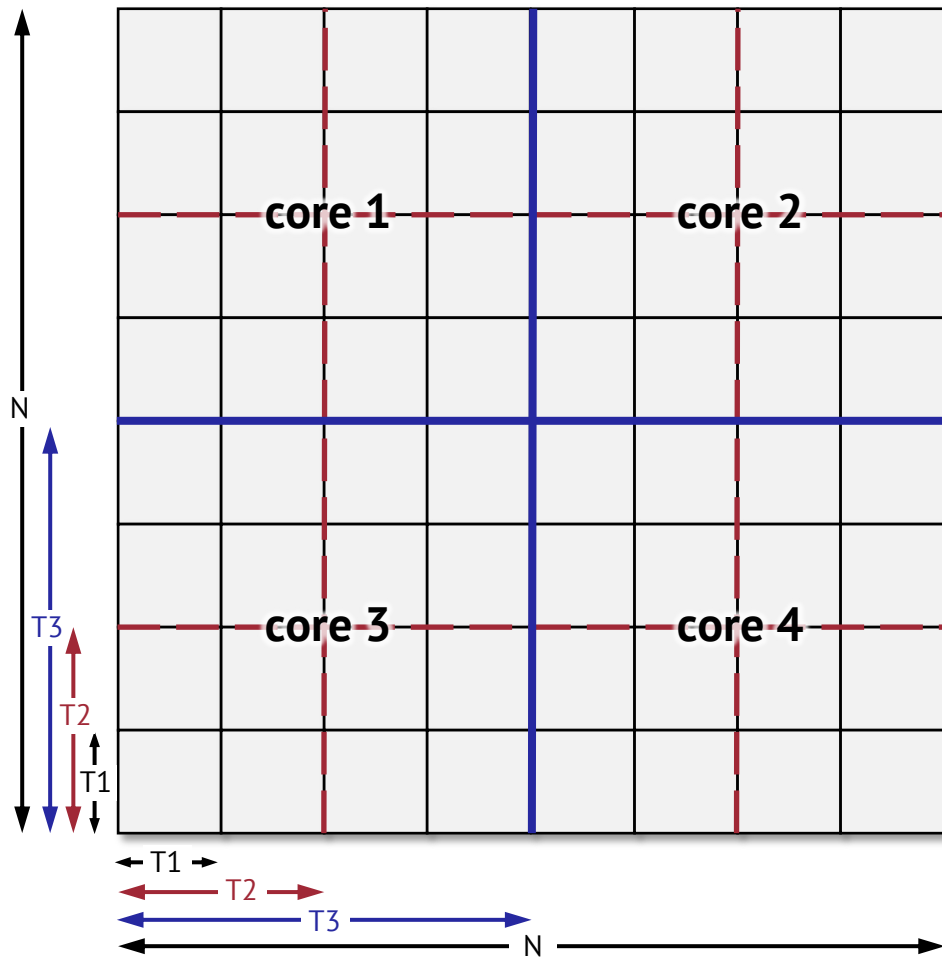


multicore reuse

- typical hierarchy
 - L1, L2 **private**
 - L3 (aka LLC) **shared**

Q. how to parallelize?

- v1: no L3 reuse T_T




multicore reuse

- typical hierarchy
 - L1, L2 **private**
 - L3 (aka LLC) **shared**

Q. how to parallelize?

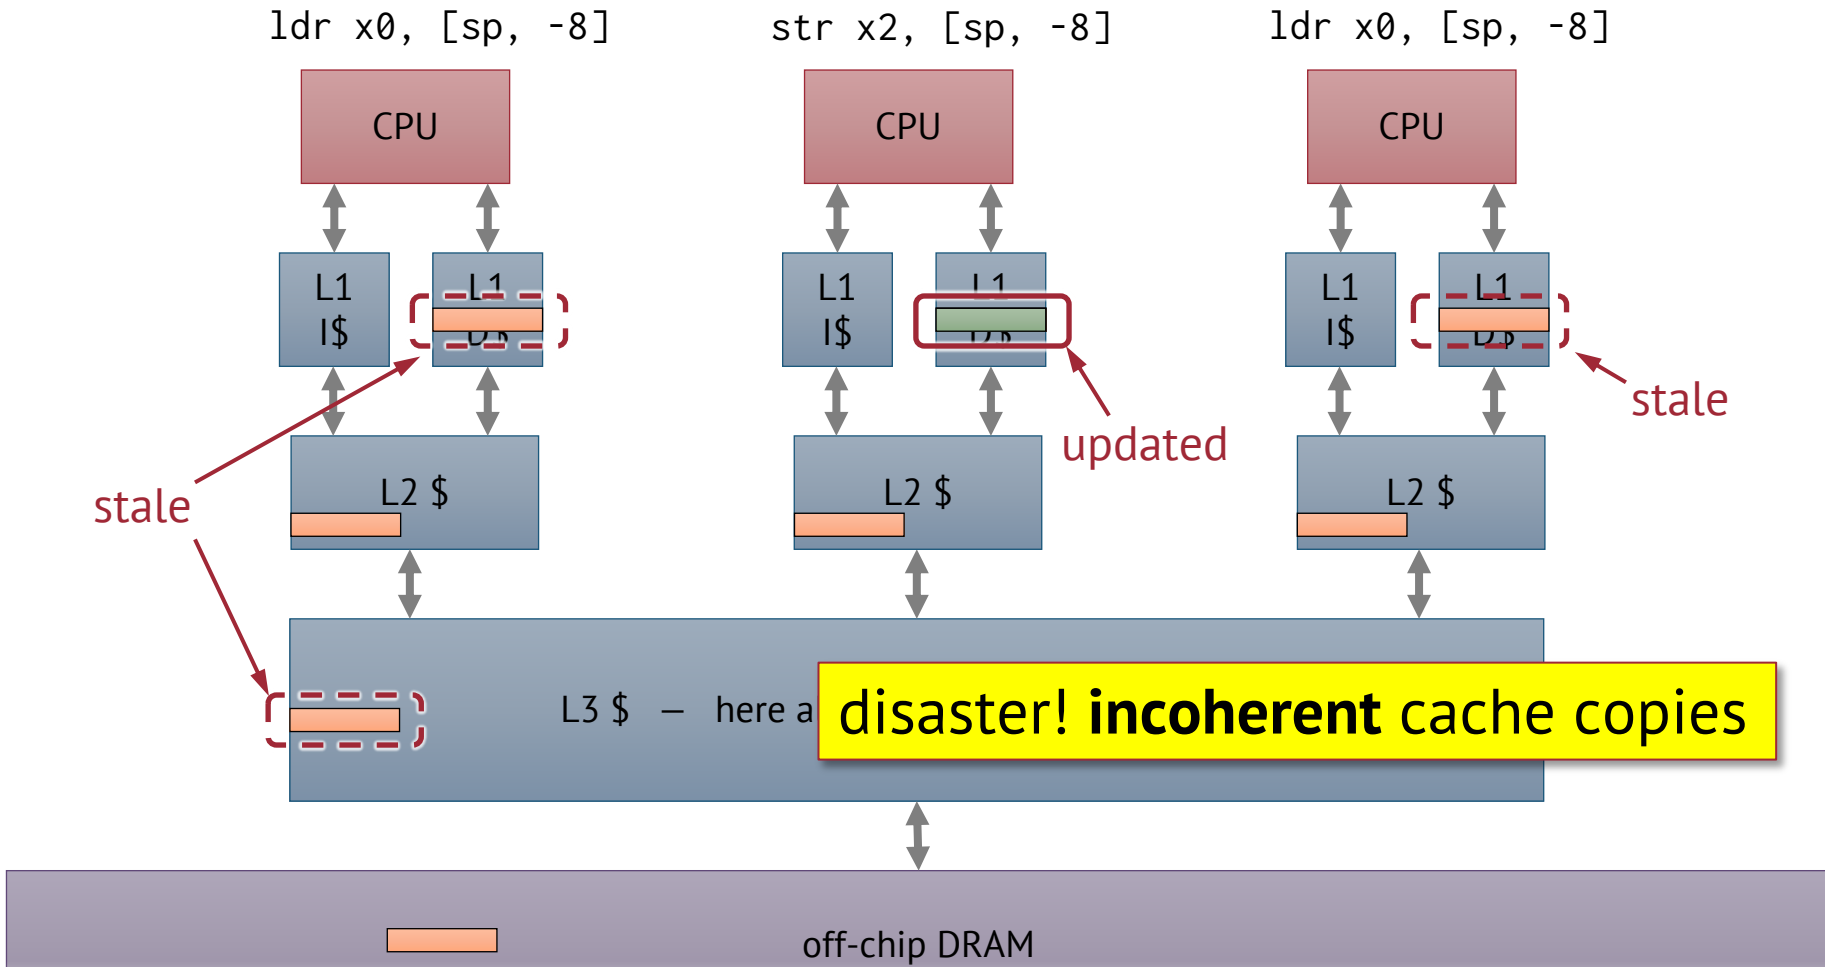
~~• v1: no L3 reuse T₁T₂~~

• v2: L3 reuse 👍



cache-aware optimization summary

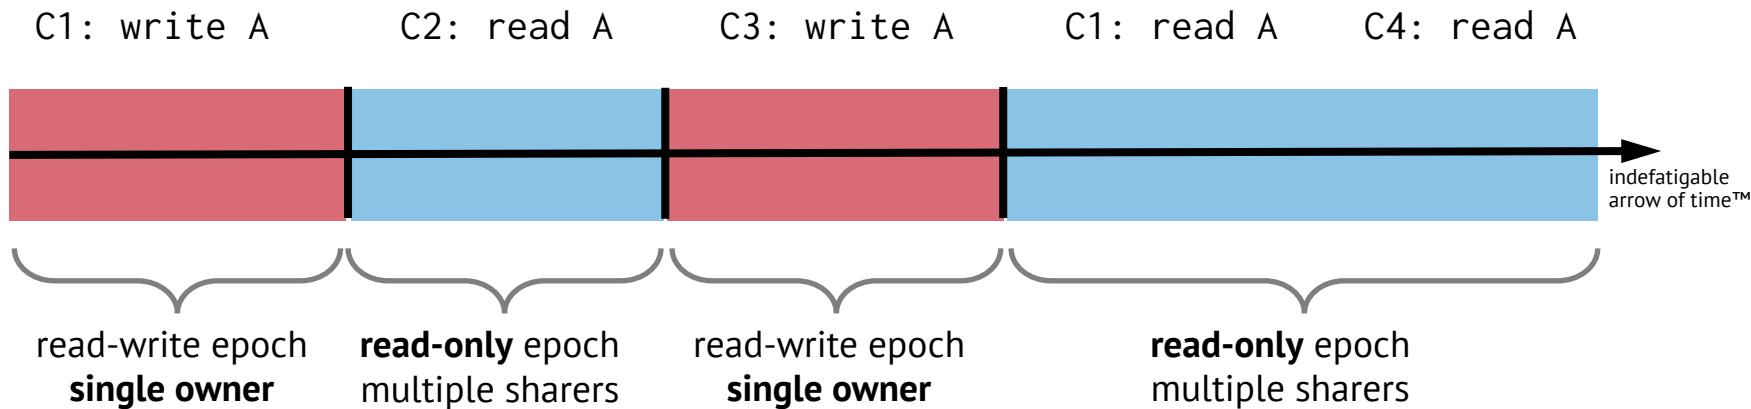
- spatial locality → **change loop bounds**
 - columns inner loop, rows outer loop (if row-major)
- temporal locality → **tile matrix**
 - size tiles so that cache holds one tile from each matrix
- cache hierarchy → **nested tiling**
 - L1\$ + L2\$ → 2 tiling levels
 - L1\$ + L2\$ + L3\$ → 3 tiling levels
 - on multicore, ensure **shared-level cache** reuse




idea: **before writing, become** the only owner




invariant: **single writer, multiple readers**





ldr x0, [sp, -8]



str x2, [sp, -8]



ldr x0, [sp, -8]



off-chip DRAM

considerations for coherent caches

- if two cores modify the same data, it will **ping-pong** between the cores

- true sharing**: same block, **same** bytes




- C1 reads cache block to access word 3
- C2 writes word 3, invalidates C1's copy of block

- false sharing**: same block, **different** bytes



- C1 reads cache block to access word 3
- C2 writes word 5, invalidates C1's copy of block *only because it's only same cache line!*



→ align R/W shared data on cache-line boundaries